



Published in HashmapInc



Senthil Anandarangan

Follow

Aug 14, 2020 · 8 min read · Listen



Save



Using Starburst Presto to Federate SQL Queries Across Multiple Data Sources

by Senthil Anandarangan and Ganesh Hegde



The data virtualization world as we know it is full of ifs and buts. Many organizations are constantly experimenting on their data management architecture and developing scalable solutions to custom fit their existing and future needs.

Although doing ETL (or ELT) to centralize large data for reporting and analytics is the preferred approach for OLAP queries, it's not always the right approach that fits all business needs. Allowing real-time access to multiple source systems' data by hiding the technical details and physical location of data and combining them into one single federated SQL query view is the challenge that Data Virtualization tools take on!

In this blog post, we want to share our experiences with trying out data virtualization capabilities of the Starburst Enterprise Presto SQL query engine querying directly on the source data of **multiple** data sources.

From One to Many

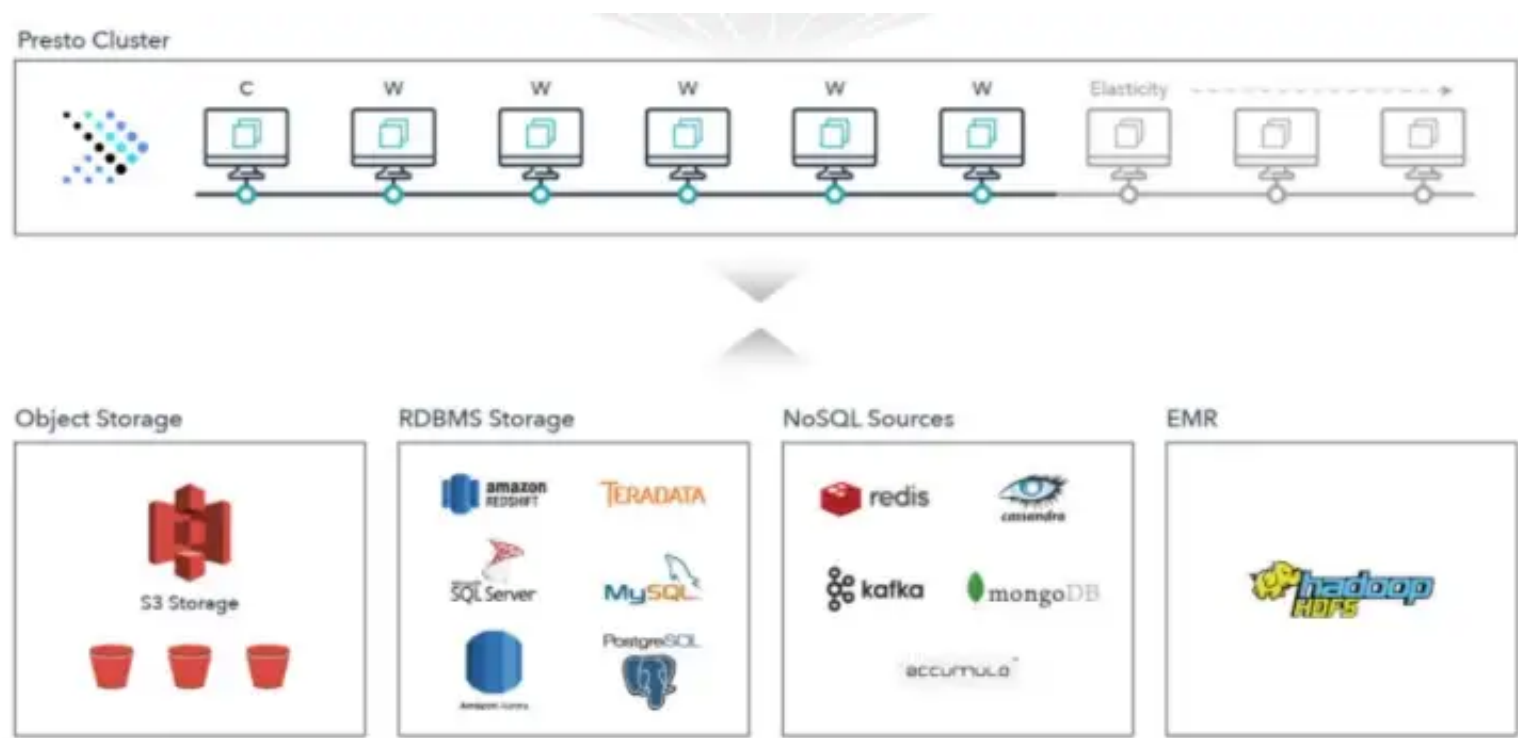
In our last blog post, we covered how to set up and run Starburst Enterprise Presto to run SQL queries directly over S3 data sources as well as how to connect Starburst Presto with Looker to create quick visualizations over the source data and analyze data patterns. Here it is in case you missed it:

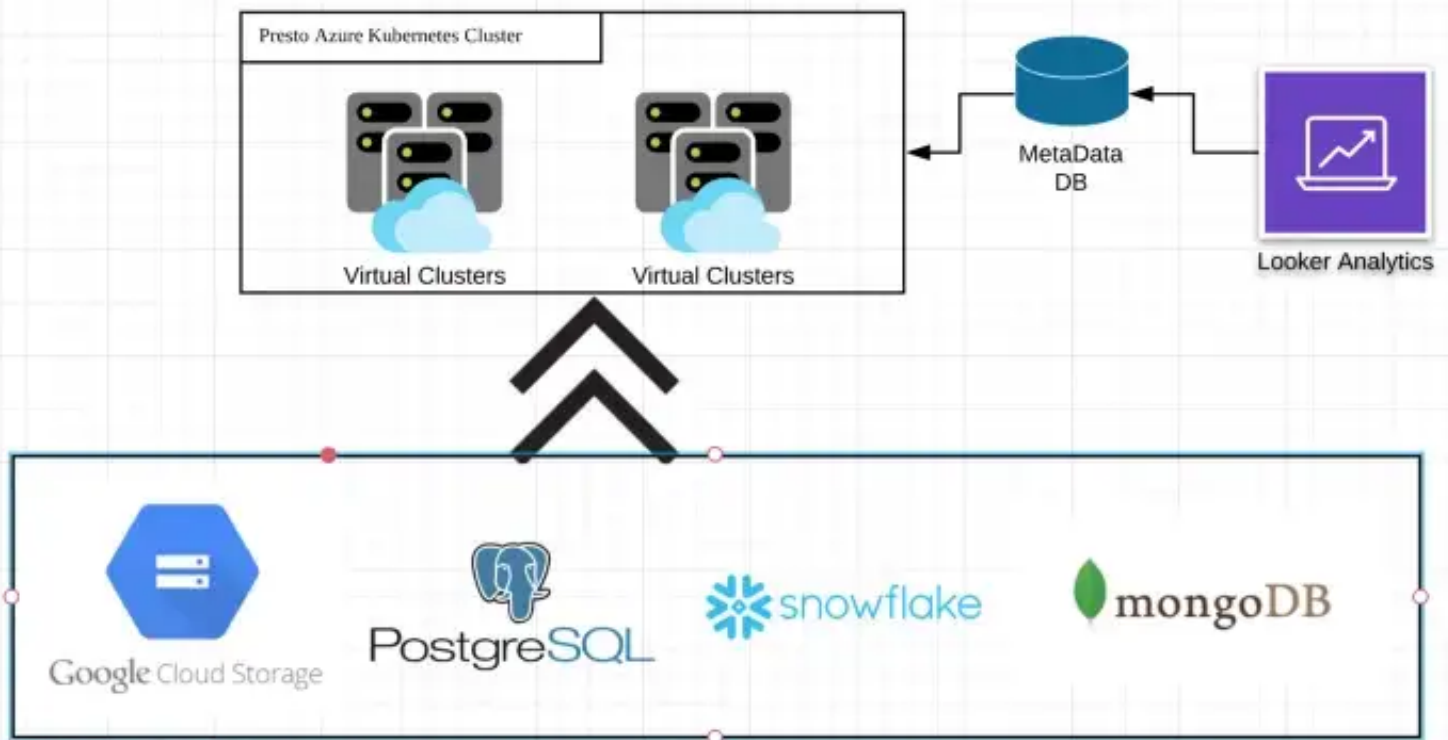
How to Connect S3 Data to Looker using Starburst Presto

In this era of increasing adoption to Data-driven solutions, faster Data accessibility becomes more and more critical...

medium.com

A single Presto query can combine data from multiple sources. Presto offers connectors to data sources including files in HDFS, AWS S3, Azure Blob/ADLS, Google Cloud Storage, MySQL, PostgreSQL, SQLServer, Oracle, AWS Redshift, Snowflake, BigQuery, Cassandra, MongoDB, Redis and many more. Refer to the Presto documentation on various data sources and connectors [here](#).





From the list of available data connectors, we picked four, created a Looker analytic view of New York subway stations, and plotted the most used stations for city bike trips.

Source Type	Source Location	Data Description
Object storage	Google Cloud Storage	NY city bike trip data
RDBMS	PostgreSQL on Azure	NY transit locations
Data Warehouse	Snowflake	NY subway stations table
No-SQL	mongoDB	NY subway station services table

Test Data source JSON:

<http://gbfs.citibikenyc.com/gbfs/gbfs.json>

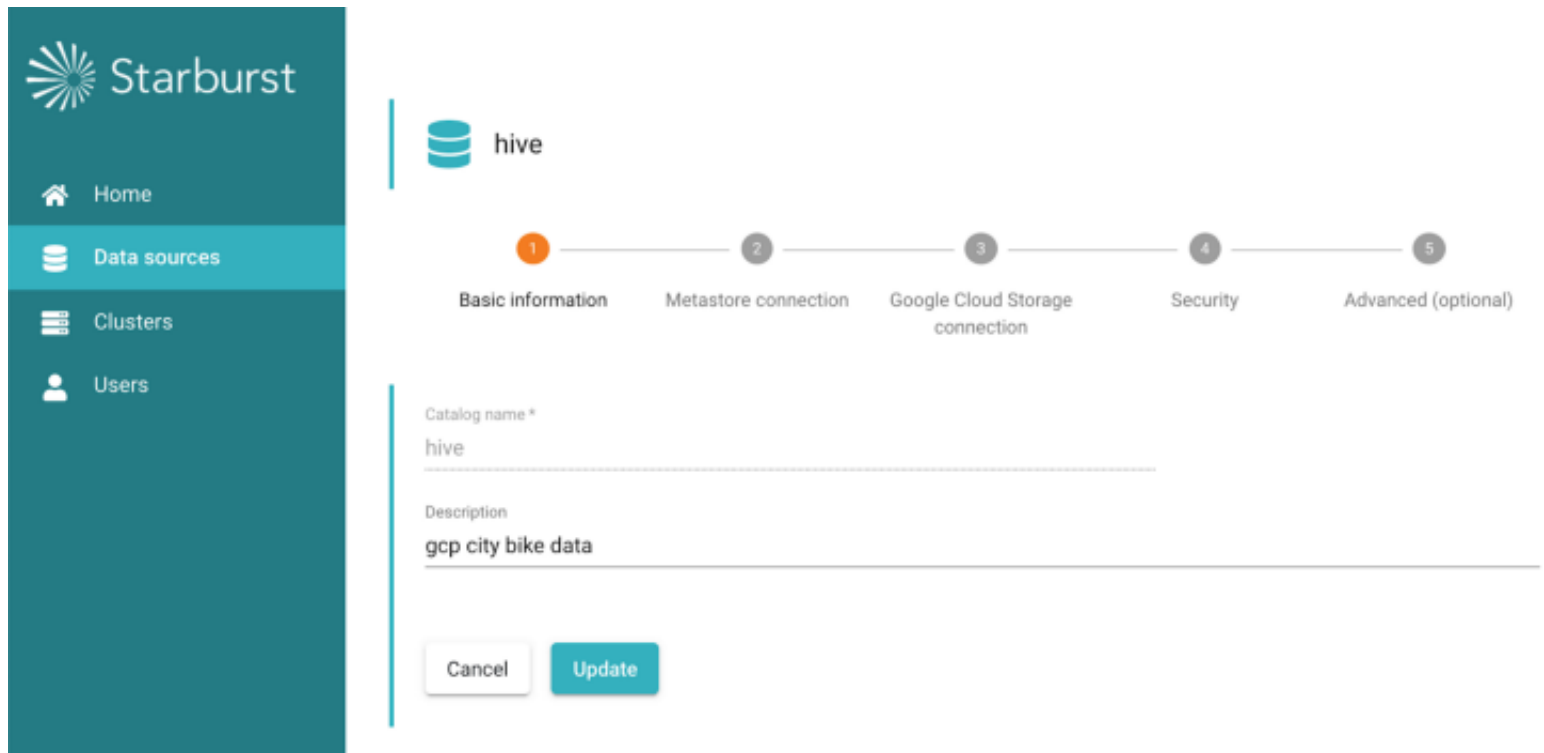
<https://www.citibikenyc.com/system-data>

Assuming, you have a running Presto Cluster setup (as described in our last [blog post](#)).


Start connecting your data sources using Mission Control UI


Steps to Connect Google Cloud Storage


1. Add + Data sources in the Mission Control and choose Google Cloud storage.
2. Note the catalog name is defaulted to 'hive' as it is enforced for object storage connections using Internal Metastore. Although it is a little confusing, this is the way Presto understands to differentiate Metastore connections in case of multiple object-store queries.
3. Choose Internal Metastore configuration and provide all details as shown.
4. Attach Google storage access key JSON for the storage location and secure storage access.





The image shows a screenshot of the Starburst web interface. On the left is a dark teal sidebar with the Starburst logo and navigation links: Home, Data sources (highlighted), Clusters, and Users. The main content area has a light gray header with a 'hive' icon and name. Below this is a horizontal progress bar with five steps: 1. Basic information (active), 2. Metastore connection, 3. Google Cloud Storage connection, 4. Security, and 5. Advanced (optional). The 'Basic information' section contains two text input fields: 'Catalog name *' with the value 'hive' and 'Description' with the value 'gcp city bike data'. At the bottom of this section are two buttons: 'Cancel' and 'Update'.


 Starburst

 Home

 Data sources

 Clusters

 Users

 hive

1

2

3

4

5

Basic informationMetastore connectionGoogle Cloud Storage connectionSecurityAdvanced (optional)

Catalog name *

hive

Description

gcp city bike data

CancelUpdate

Metastore location *

☐ External Hive Metastore
☒ Internal Metastore

Internal Metastore database *

Internal Metastore with PostgreSQL RDBMS

Database host *

starburst-hive.postgres.database.azure.com

Database port *

5432

Database name *

hivemetastore

JDBC connection URL

jdbc:postgresql://starburst-hive.postgres.database.azure.com:5432/hivemetastore

Database user *

postgres@starburst-hive

Database password *

Cancel

Update



hive


Service account key file *

Drop *.json file here or click to upload



senthil-anandarangan-poc-301fe62a363d.json
2.31 KB


Cancel

Update



Home
Data sources
Clusters
Users



Basic information
Metastore connection
Google Cloud Storage connection
Security
Advanced (optional)

Write options

Default storage format *
ORC

File compression codec *
Gzip

☐ Immutable partitions

☒ Enable writing to external Hive tables

☒ Collect column level statistics on write

Read options

Maximum split size *
64 MB

Maximum initial splits count *
200

☐ Force local scheduling

Advanced configuration

☐ Manually configure catalog properties

Cancel Update

☒ Respect existing tables format

Maximum number of partitions per write *
100

☐ Create empty files for buckets without data

☒ Enable creating external Hive tables

Maximum number of partitions for a table scan *
100000

Maximum initial split size *
32 MB

☐ Recursive directories

After the update, open cluster configuration in Mission Control and add/include Google Cloud storage data source.

That's it! We have configured the Google Cloud storage dataset. In order to start querying, it requires a restart of the Mission Control Clusters, however, we will hold on to restart until we connect the other three data sources.

Steps to Connect PostgreSQL on Azure Cloud

1. Add + Data sources in the Mission Control and choose PostgreSQL connector.
2. The configuration is very straight forward here; only the DB connection string parameters that you would provide here to establish a connection.

Azure services



Azure Database
for PostgreS...

Steps to Connect to MongoDB no-SQL database

1. Add + Data sources in the Mission Control and choose MongoDB connector.
2. Provide the Mongo seed string to establish a connection.

Steps to Connect to Snowflake Data Warehouse

Although the Snowflake connector is available and supported by Presto, my Starburst version is not the latest and it was not listed in the available data sources to Add. There are 2 alternate ways to add unlisted connectors in Mission Control (of course, only if they are supported by Presto!).

1. Check if you find a link on the data sources screen that says 'Add one manually'. Once you add a manual data source, it will look like the example below in Snowflake:



The screenshot shows a configuration form for a Snowflake connector. At the top left is the Snowflake logo. Below it, the 'Catalog name' field is filled with 'snowflake'. The 'Description' field is empty. The 'Connector name' field is filled with 'snowflake-jdbc'. Below these fields is a warning message: 'Incorrect configuration may result in cluster startup failures or problem with the catalog usage.' At the bottom, there is a text area for 'Catalog properties (properties file format)' containing the following text: 'connection-url=jdbc:snowflake://gh41405.snowflakecomputing.com/', 'connection-password=[REDACTED]', 'connection-user=[REDACTED]', 'snowflake.warehouse=Demo', 'snowflake.database=DEMO_DB', and 'snowflake.role=SYSADMIN'. At the bottom left is a 'Cancel' button, and at the bottom right is an 'Update' button.

 snowflake

Catalog name *

snowflake

Description

Connector name *

snowflake-jdbc

 Incorrect configuration may result in cluster startup failures or problem with the catalog usage.

Catalog properties (properties file format)

```
connection-url=jdbc:snowflake://gh41405.snowflakecomputing.com/
connection-password=[REDACTED]
connection-user=[REDACTED]
snowflake.warehouse=Demo
snowflake.database=DEMO_DB
snowflake.role=SYSADMIN
```

Cancel Update

2. If you do not see a Manual Add option, then you are probably not on the latest release of Mission Control. The alternative is to get .22 version installed where you can create manual

connectors. Since we do not want to upgrade at this point, here is the other workaround that added Snowflake as bootstrap entry:

To add a connector that isn't supported by Mission Control, you will need to create a bootstrap file that adds this file to the `/etc/Presto/catalog` directory of the Co-ordinator POD:

```
senthilkumar@Azure:~$ kubectl get pods
```

NAME	READY	STATUS
hive-metastore-mc-7c40400c29a4430-5f976f7dc9-8q4nw	1/1	Running
mission-control-6dd7c845c5-2tqvn	1/1	Running
postgres-0	1/1	Running
presto-coordinator-mc-7c40400c29a4430-d4ffc48fd-x6wmh	2/2	Running
presto-operator-75f8dd9f8c-ncxq8	1/1	Running
presto-worker-mc-7c40400c29a4430-d7cc567c8-m2k5b	1/1	Running

Using a bootstrap:

1. Create yaml file below (shown as an example, any connector would work in here):

```
$ more snowflake.yaml
```

```
kind: ConfigMap
```

```
apiVersion: v1
```

```
metadata:
```

```
name: snowflake-connector
```

```
data:
```

```
bootstrap.sh: |-
```

```
tee << EOF > /etc/presto/catalog/snowflake.properties
```

```
connector.name=snowflake-jdbc
```

```
connection-url=jdbc:snowflake://xxxxxxxxxxxxx.snowflakecomputing.com/
```

```
connection-password=xxxxxxxxxxxxxxxxxxxx
```

```
connection-user=SNOWFLAKE_DEMO
```

```
snowflake.warehouse=Demo
```

```
snowflake.database=DEMO_DB
```

```
snowflake.role=SYSADMIN
```

```
EOF
```

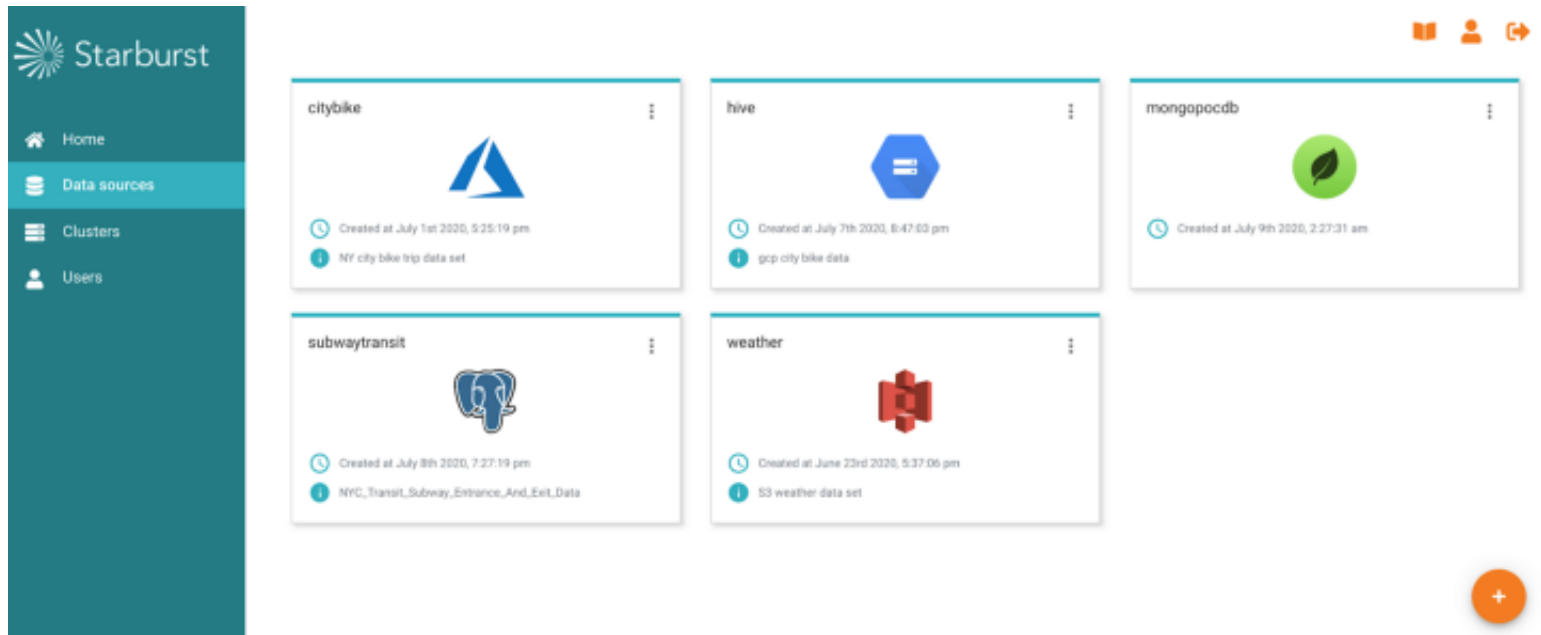
2. `kubectl apply -f snowflake.yaml`

3. Under the General Configuration click on the “Add Bootstrap script volume” and fill “snowflake-connector” as the name of the ConfigMap.

4. Save and restart the cluster.

Multiple Data Sources Connected View

After you connected all the above data sources, the view should look like below:



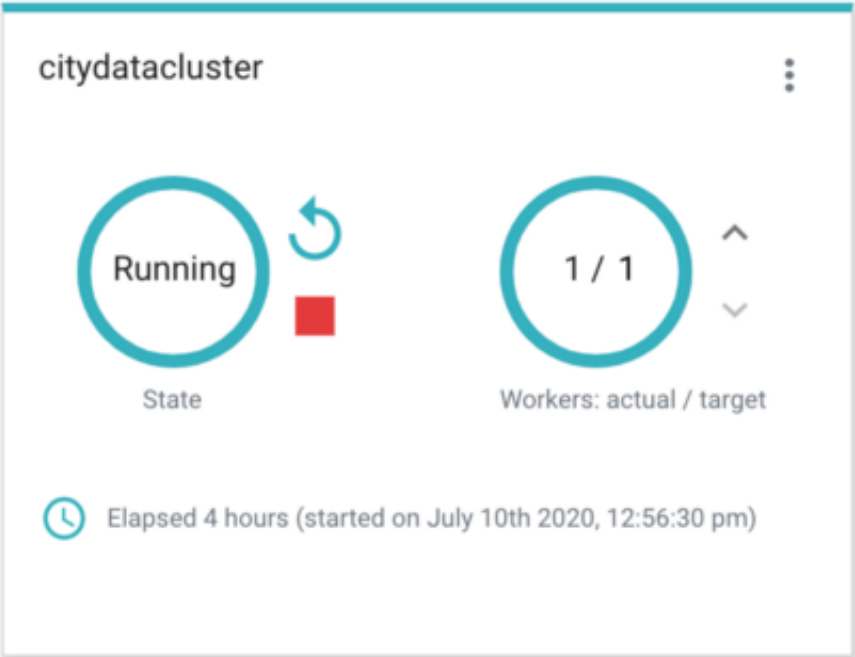

There are a few important things to note here:

1. Azure ADLS and AWS S3 are also present from my other project and are not used for this demo. We can have as many data sources connected and can selectively configure [add/remove] them in the Run Cluster configuration.
2. Keep in mind, Presto mandates a hive metastore connection for object storage connections. You can configure a connection to multiple object stores (ADLS, S3, etc.) and use the same Presto cluster to query them. For the 2nd connection (ADLS), configure it to use an external hive metastore and point it to the same one being used by the 1st (S3) connection. In terms of configuration, suppose if you have S3 and ADLS and are trying to query from both in the same SQL query, then the S3 connection will have Internal Metastore configured to the Hive Metastore DB connection parameters. ADLS will have the External Hive Metastore configured and pointing to the Internal Metastore connection hostname which should look like 'hive-metastore-mc-0xxxxxxx.default.svc.cluster.local' (if you are in doubt, this hostname can be located in the hive.properties file in /etc/presto/catalog of Co-ordinator POD).
3. Snowflake is missing in the list. This is because we have manually configured the Snowflake connector as it is not automatically available in my version of Starburst Demo. Though it doesn't show up here, Snowflake tables can be still be queried (from Looker or another tool).

If you want to double-check this type of connector installation, you may do so by looking at the Co-ordinator POD directory. The '/etc/presto/catalog' directory should list all connected data sources. Look for 'snowflake.properties' here, like for any other source .properties file.

```
senthilkumar@Azure:~$ kubectl exec -it presto-coordinator-mc-7c40400c29a4430-d4ffc48fd-x6wmh sh
Defaulting container name to presto-coordinator.
Use 'kubectl describe pod/presto-coordinator-mc-7c40400c29a4430-d4ffc48fd-x6wmh -n default' to see all of the containers in this pod.
sh-4.2# ls /etc/presto/catalog/
hive.properties  mongopocdb.properties  snowflake.properties  subwaytransit.properties  tpch.properties  weather.properties
sh-4.2#
```

It's time to Restart and Run the Presto Cluster!



Verify the cluster shows 'Running' state.

Connecting Looker to Starburst Presto and Multiple Databases

Refer to our previous [blog post](#) outlining details on connecting Looker to Presto.

How to Connect S3 Data to Looker using Starburst Presto

In this era of increasing adoption to Data-driven solutions, faster Data accessibility becomes more and more critical...

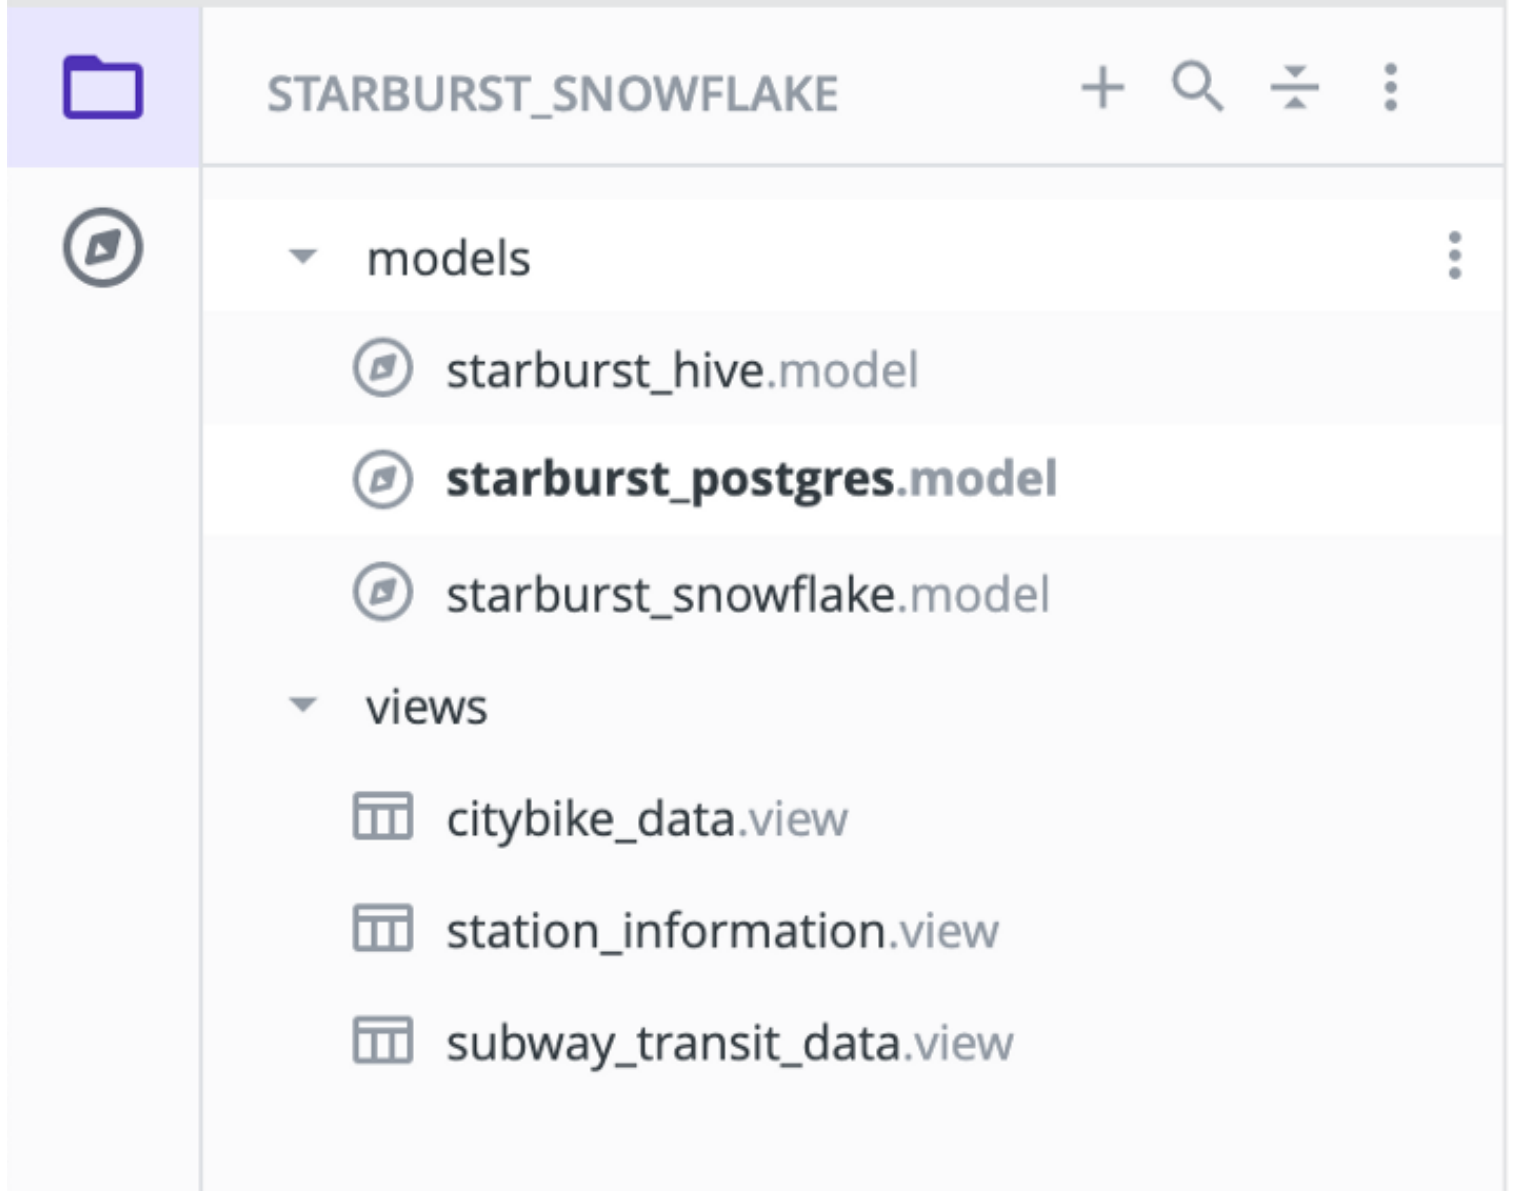
medium.com

Setup new connections as stated below to connect to multiple data sources. Note, ‘hive’ is the mandatory data source name for the object sources using Internal MetaStore connections. More detail on this can be found [here](#).

● starburst_multi	hive
● starburst_snowflake	snowflake
● starbust_poc	hive
● starbust_postgres	subwaytransit

Developing LookML

Under “Develop -> (your Project)” you should see the tables appear for writing LookML Models and Views (one model per database):



Joining tables across connections

In the model, join the views with primary and foreign keys. If no such relationship exists, join by common field name.

```
14 | join: subway_transit_data |  
15 |   sql_on: station_id = CAST(citybike_data.startstationid AS DECIMAL(10, 4)) ;;  
16 |   relationship: one_to_many  
17 | }
```

Exploring the data in Looker

Explore the data before creating a dashboard by selecting multiple dimensions and measures with sorting and filters.

Explore

Station Information

Search

All Fields

Dimensions

Measures

▶ Citybike Data

1

▶ Station Information

2

▶ Subway Transit Data

Citybike Data

Dimensions

- Bikeid
- Birthyear
- Endstationlatitude
- Endstationlongitude
- Endstationname
- Gender
- Startstationid **PIVOT** **FILTER**
- Startstationlatitude
- Startstationlongitude
- Startstationname
- Starttime
- Stoptime
- Tripduration
- Usertype

Measures

- Count **FILTER**

Station Information

VISUALIZATION

Startstationid

Citybike Data

1	3186	4,499
2	3203	2,385
3	3639	1,830
4	3202	1,782
5	3195	1,740
6	3276	1,476
7	3184	1,400
8	3199	1,276

DATA **RESULTS** **SQL** **Calculations** **Row Limit** 500 **Totals**


Citybike Data Startstationid	Citybike Data Count
1 3186	4,499
2 3203	2,385
3 3639	1,830
4 3202	1,782
5 3195	1,740
6 3276	1,476
7 3184	1,400
8 3199	1,276

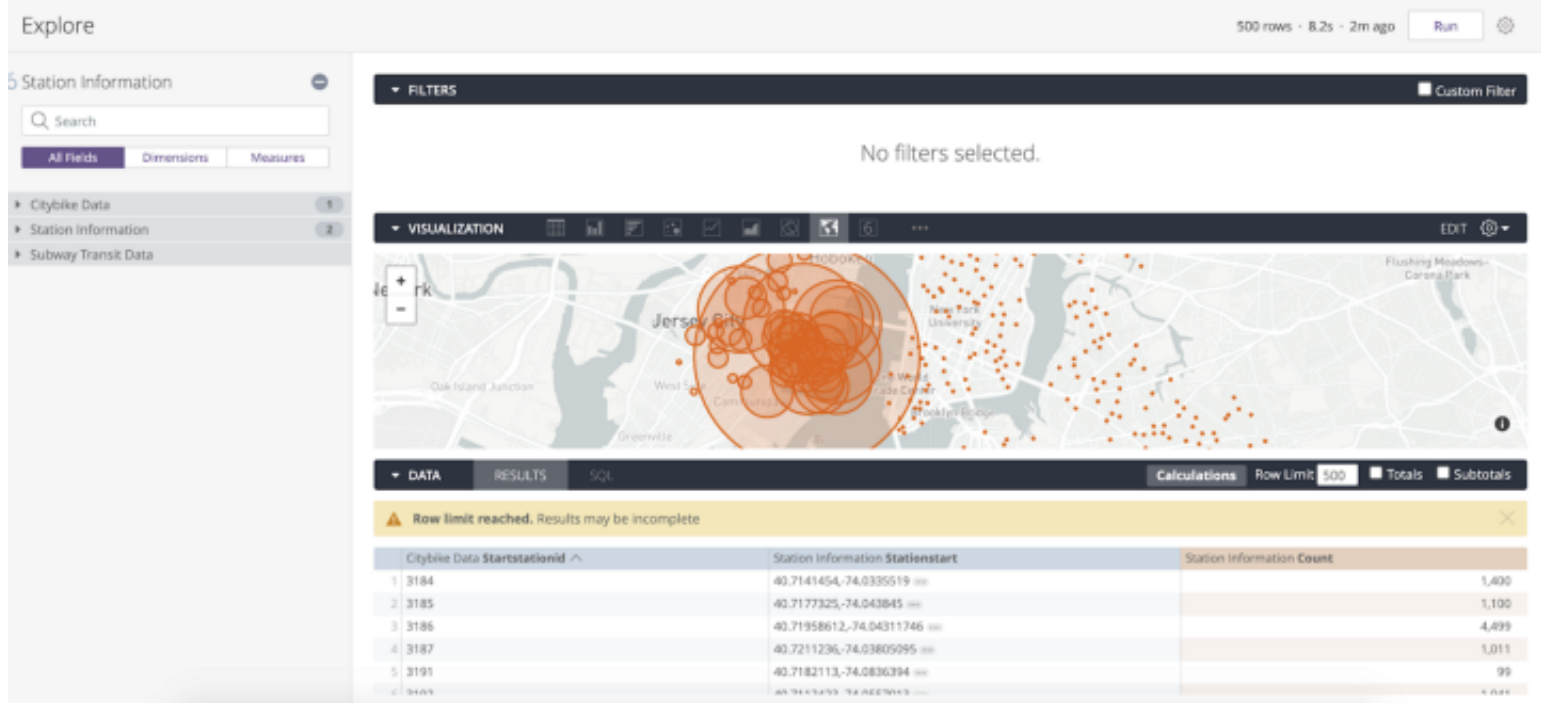
Open in app ↗

Get unlimited access

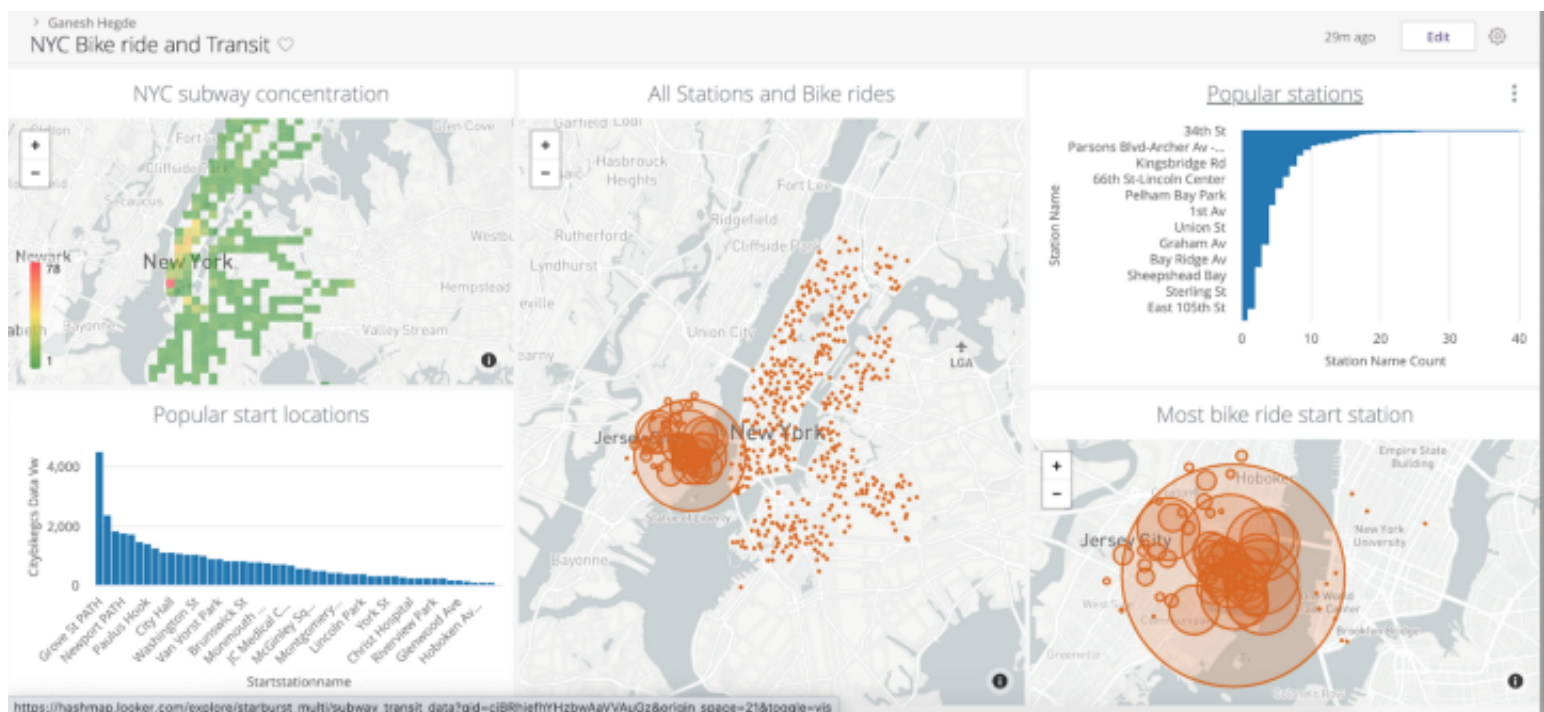


Search Medium





This dashboard is built using city bike data from Hive (Google Cloud storage), station data from Snowflake, and transit data from Postgres with Query federation done using Starburst Presto.



Closing Thoughts

With the growing list of data connectors, Presto provides an opportunity to realize data virtualization with federated SQL queries across multiple data sources. Presto's Query Plan viewer is a huge plus for developers to identify SQL execution plans as well as optimize and improve query response times. Share your experiences with us if you have tried connecting other data sources and reporting tools to Starburst Presto!

Need Help with Your Cloud Initiatives?

If you are considering migrating or modernizing data and analytics products and applications in the cloud or if you would like help and guidance and a few best practices in delivering higher value outcomes in your existing cloud program, then please [contact us](#).

[Hashmap](#) offers a range of enablement workshops and assessment services, cloud modernization and migration services, and consulting service packages as part of our service offerings.

Feel free to share on other channels and be sure and keep up with all new content from Hashmap [here](#). To listen in on a casual conversation about all things data engineering and the cloud, check out Hashmap's podcast [Hashmap on Tap](#) as well on Spotify, Apple, Google, and other popular streaming apps.

<div><div>Hashmap on Tap Hashmap Podcast</div><div>A rotating cast of Hashmap hosts and special guests explore technologies from diverse perspectives while enjoying a drink of choice.</div></div>	
--	--

[Senthil Anandarangan](#) and Ganesh Hedge are Cloud, Data, and UI Engineers with [Hashmap](#) providing Data, Cloud, IoT, and AI/ML solutions and consulting expertise across industries with a group of innovative technologists and domain experts accelerating high-value business outcomes for our customers.