



Geoffrey Koh

Follow

Apr 4, 2020 · 6 min read · [Listen](#)



Save



Fun with Fixtures for Database Applications

Integration testing with databases is one of the most vital, yet commonly overlooked part of any software development process. This is, in part, due to the need to have a reachable database server active around somewhere. In this article, we explore unit and integration testing with databases that can be repeatedly spun up for testing, provides a uniform schema and test data, and can be easily torn down afterwards. All these can be done with the help of some nifty open source Python libraries that are readily available.



A Little Bit on Pytest

`pytest` is the current *de facto* testing framework for any Python application, surpassing in popularity to the `unittest` module, which comes as a standard package with all modern Python distributions. Vanilla use of `pytest` is very straight-forward. All one needs to do is to write a function with `test` as a substring of the function name, and run the `pytest` command on it.

```
1 def add(a, b):
2     """ Simple function to add two numbers """
3     return a + b
4
5 def test_add():
6     assert add(2, 4) == 6, 'The correct answer should be 6'
```

medium_pytest_simple.py hosted with ♥ by GitHub

[view raw](#)

```
(base) Geoffreys-MBP:pytest geoffkoh$ pytest medium_pytest_simple.py
===== test session starts =====
platform darwin -- Python 3.7.0, pytest-3.8.0, py-1.6.0, pluggy-0.7.1
rootdir: /Users/geoffkoh/Development/Python/medium/pytest, inifile:
plugins: xdist-1.26.1, remotedata-0.3.0, openfiles-0.3.0, forked-1.1.2, doctestplus-0.1.3, cov-2.8.1, arraydiff-0.2
collected 1 item

medium_pytest_simple.py . [100%]

===== 1 passed in 0.03 seconds =====
```

Example of testing a simple function with `pytest`

However, the real power of `pytest` comes when one understands all the bells and whistles underneath the hood, in particular **fixtures**, which can help perform various setting up and tearing down of relevant environments required to run test cases.

There are a lot of very good tutorials on `pytest` and thus we will not go into them in detail. This article focuses purely on using fixtures to set up temporary databases to achieve a more realistic and repeatable testing environment.

Fixtures

Fixtures are objects that set up certain conditions that are then used in testing. Think of them as functions that are called before running the actual test functions. In `pytest`, this can be done by marking a function with the decorator `@pytest.fixture`. Thereafter, you can use the fixture in any test function, simply by adding it as an argument to the function definition. A simple example is shown below.

```

1 # Standard imports
2 import requests
3 import sqlite3
4
5 # Third party imports
6 import pytest
7
8 @pytest.fixture
9 @pytest.fixture
10 def setup_database():
11     """ Fixture to set up the in-memory database with test data """
12     conn = sqlite3.connect(':memory:')
13     cursor = conn.cursor()
14     cursor.execute('''
15         CREATE TABLE stocks
16         (date text, trans text, symbol text, qty real, price real)''')
17     yield conn
18
19 @pytest.fixture
20 def setup_test_data1(setup_database):
21     cursor = setup_database
22     sample_data = [
23         ('2020-01-01', 'BUY', 'IBM', 1000, 45.0),
24         ('2020-01-01', 'SELL', 'GOOG', 40, 123.0),
25     ]
26     cursor.executemany('INSERT INTO stocks VALUES(?, ?, ?, ?, ?)', sample_data)
27     yield cursor
28
29 @pytest.fixture
30 def setup_test_data2(setup_database):
31     cursor = setup_database
32     sample_data = [
33         ('2020-01-01', 'SELL', 'TESLA', 400, 233.0),
34         ('2020-01-01', 'SELL', 'MSFT', 140, 980.0),
35         ('2020-02-01', 'BUY', 'AMZN', 3000, 1200.0),
36     ]
37     cursor.executemany('INSERT INTO stocks VALUES(?, ?, ?, ?, ?)', sample_data)
38     yield cursor
39
40 def test_with_sample_data2(setup_test_data2):
41     # Test to make sure that there are 3 items in the database
42     cursor = setup_test_data2
43     assert list(cursor.execute('SELECT * FROM stocks')) == 3

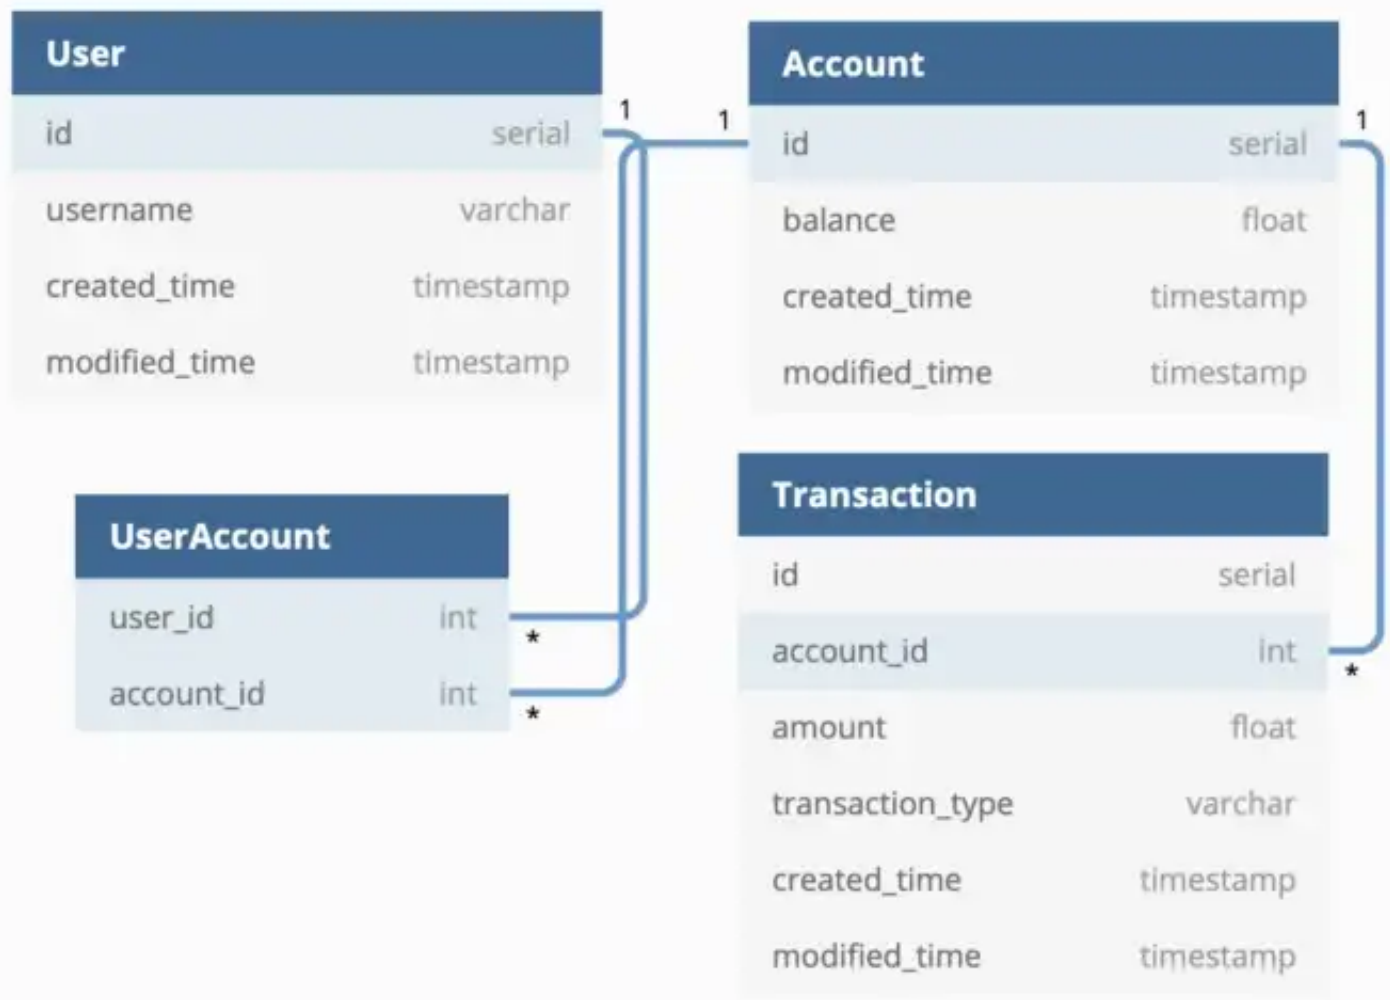
```

Later we will see some of these being used in the test codes for database connections.

Testing Database Interactions with Fixtures

We shall go through this tutorial with a simple example of an bank account transaction system (I mentioned in my previous article that I will not use examples that have no semblance to real life applications, but yet at the same time, simple enough to illustrate the concepts). As a proponent of **good clean design**, I will also recommend using `SQLAlchemy` (<https://www.sqlalchemy.org>) to model the various business objects.

In this system, each customer will be a **user**, and each user can have one or more **accounts**. At the same time, an **account** can belong to a single user, or it can be a joint account belonging to **multiple users**. Every time the user credits or debits into the account, a **transaction** will be created. The Entity-Relation diagram reflecting the relationship is shown below.



ER Diagram for the simple bank account application

Now that we have the business objects, their Object Relation Model definitions in SQLAlchemy is then shown in the snippet below.

```

1  # Third party imports
2  from sqlalchemy import Column, Integer, String, DateTime, Float, ForeignKey
3  from sqlalchemy.ext.declarative import declarative_base
4  from sqlalchemy.orm import relationship
5  from sqlalchemy.sql import func
6
7  Base = declarative_base()
8
9
10 class User(Base):
11     """ Model class to represent a user """
12
13     __tablename__ = 'user'
14
15     id = Column(Integer, primary_key=True, autoincrement=True)
16     username = Column(String, nullable=False, unique=True)

```

```
17         created_time = Column(DateTime(timezone=True),
18                                 server_default=func.now())
19         modified_time = Column(DateTime(timezone=True),
20                                 server_default=func.now(),
21                                 onupdate=func.now())
22
```

```
1  # Standard imports
2  from typing import List
3
4  # Third party imports
5  from sqlalchemy.orm.session import Session
6  from sqlalchemy.orm.exc import NoResultFound
7
8  # Application imports
9  from models import Account, User, Transaction
10
11
12 def get_user(username: str, session: Session) -> User:
13     """ Gets the user by username """
14
15     try:
16         user = session.query(User).filter(User.username == username).one()
17         return user
18     except NoResultFound:
19         return None
20
21 # end get_user()
22
23
24 def get_accounts_by_user(username: str, session: Session) -> List[Account]:
25     """ Retrieves the accounts given the username """
26
27     user = get_user(username, session)
28     accounts = list(session.query(Account).filter(
29         Account.users.contains(user)).all())
30     return accounts
31
32 # end get_accounts_by_user
33
34
35 def compute_balance(username: str, session: Session) -> float:
36     """ Computes the balance based on the username """
37
38     accounts = get_accounts_by_user(username, session)
39     balance = sum([account.balance for account in accounts])
40     return balance
41
42 # end compute_balance()
```

```

43
44
45 def debit(account_id: int, amount: float, session: Session) -> float:
46     """ Debits amount into the account """
47
48     try:
49         account = session.query(Account).filter(Account.id == account_id).one()
50         account.balance -= amount
51         transaction = Transaction(account_id=account_id,
52                                   amount=amount,
53                                   transaction_type='DEBIT')
54         session.add(transaction)
55         session.commit()
56         return account.balance
57     except NoResultFound:
58         return 0
59
60 # end debit()

```

```

61
1  # Third party imports
2  import pytest
3  from sqlalchemy import create_engine
4  from sqlalchemy.orm import sessionmaker
5
6  # Application imports
7  from core import get_user, get_accounts_by_user, compute_balance, debit
8  from models import Base, User, Account, UserAccount, Transaction
9
10
11 @pytest.fixture(scope='function')
12 def setup_database():
13
14     engine = create_engine('sqlite://')
15     Base.metadata.create_all(engine)
16     Session = sessionmaker(bind=engine)
17     session = Session()
18     yield session
19     session.close()
20
21 # end setup_database()
22
23
24 @pytest.fixture(scope='function')
25 def dataset(setup_database):
26
27     session = setup_database
28
29     # Creates user

```

```
30     john = User(username='john')
31     mary = User(username='mary')
32     session.add(john)
```

```
(base) Geoffreys-MBP:pytest geoffkoh$ pytest -sv medium_pytest_functions.py
===== test session starts =====
platform darwin -- Python 3.7.0, pytest-3.8.0, py-1.6.0, pluggy-0.7.1 -- /anaconda3/bin/python
cachedir: .pytest_cache
rootdir: /Users/geoffkoh/Development/Python/medium/pytest, inifile:
plugins: xdist-1.26.1, remotedata-0.3.0, openfiles-0.3.0, forked-1.1.2, doctestplus-0.1.3, cov-2.8.1, arraydiff-0.2
collected 1 item
```

```
medium_pytest_functions.py::test_database PASSED
```

```
39     joint_account = Account(balance=20.0)
40     john.accounts.append(john_account)
41     mary.accounts.append(mary_account)
42     john.accounts.append(joint_account)
43     mary.accounts.append(joint_account)
44     session.add(john_account)
45     session.add(mary_account)
46     session.add(joint_account)
47     session.commit()
48
49     yield session
50
51 # end dataset_1
52
53
54 def test_database(dataset):
55
56     # Gets the session from the fixture
57     session = dataset
58
59     # Do some basic checking
60     assert len(session.query(User).all()) == 2
61     assert len(session.query(Account).all()) == 3
62     assert len(session.query(UserAccount).all()) == 4
63
64     # Retrieves John and Mary
65     john = get_user('john', session)
66     mary = get_user('mary', session)
67
68     # Checks their accounts
69     assert len(get_accounts_by_user(john.username, session)) == 2
70     assert len(get_accounts_by_user(mary.username, session)) == 2
71
72     # Checks the balance
73     assert compute_balance(john.username, session) == 30.0
74     assert compute_balance(mary.username, session) == 25.0
```

```
1 # Standard imports
2 import tempfile
3
```



```

4  # Third party imports
5  import pytest
6  from sqlalchemy import create_engine
7  from sqlalchemy.orm import sessionmaker
8  from pytest_postgresql import factories
9
10
11 # Using the factory to create a postgresql instance
12 socket_dir = tempfile.TemporaryDirectory()
13 postgresql_my_proc = factories.postgresql_proc(
14     port=None, unixsocketdir=socket_dir.name)
15 postgresql_my = factories.postgresql('postgresql_my_proc')
16
17
18 @pytest.fixture(scope='function')
19 def setup_database(postgresql_my):
20
21     def dbcreator():
22         return postgresql_my.cursor().connection
23
24     engine = create_engine('postgresql+psycopg2://', creator=dbcreator)
25     Base.metadata.create_all(engine)
26     Session = sessionmaker(bind=engine)
27     session = Session()
28     yield session
29     session.close()
30
31 # end setup_database()

```

medium_pytest_functions_pgsql.py hosted with ♥ by GitHub

[view raw](#)

- `pytest-redis` — <https://pypi.org/project/pytest-redis/>
- `pytest-consul` — <https://pypi.org/project/pytest-consul/>
- `pytest-kafka` — <https://pypi.org/project/pytest-kafka/>
- ... And more

Coincidentally, they are all created with similar interfaces, such that the process of setting them up goes through roughly the same processes.

Hopefully this article can give users a clearer understanding of how to better test database-driven applications, and thereby leading to more robust and reliable software.

Given the ubiquity of Python and `pytest`, if you find that a plugin for a database you want to test for is missing, you can either write one (or contact me), or appeal to the general Python